


CUDA Optimization with Parallel Transpose

이 정 근 (Jeong-Gun Lee)

한림대학교 소프트웨어융합대학, 임베디드 SoC 연구실
www.onchip.net
 Email: Jeonggun.Lee@hallym.ac.kr




Before Start ~

- How to use a **Google Drive** in a **Colab** environment ?

Google Drive를 활용한 CUDA 프로그래밍
 직접 Google Drive에서 프로그램을 수행할 수 있습니다.
 이경우 mount가 필요하고 여러개의 파일로 이루어진 코드들을 수행할 수 있습니다.
 또한 직접 구글 드라이브로 **git clone**하여 소스를 살펴보고 실행 데이터와 함께 수행 시킬 수 있습니다.

```
In [0]: from google.colab import drive
Enter your authorization code:
지정된 URL을 클릭하면 인증 코드를 볼 수 있으며, 이를 복사하여 입력합니다.

In [2]: drive.mount('/content/gdrive')
Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3Aietf%3Aawg%3Aoauth%3A2.0%3Aaob&scope=email%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdocs.test%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fdrive.photos.readonly%20https%3A%2F%2Fwww.googleapis.com%2Fauth%2Fpeopleapi.readonly&response_type=code Enter your authorization code: ..... Mounted at /content/gdrive

In [3]: %cd /content
/content

In [4]: !ls
gdrive sample_data


In [5]: %cd gdrive
/content/gdrive
```

https://github.com/jeonggunlee/CUDATeaching/blob/master/02_cuda_lab/00_googleDrive_CUDAExam.ipynb

CUDA Optimization – Matrix Transpose

- 최적화의 예로 “**Optimizing Parallel Transpose in CUDA**”를 살펴보도록 해요!

<https://devblogs.nvidia.com/efficient-matrix-transpose-cuda-cc/>

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$


Matrix Transpose 1

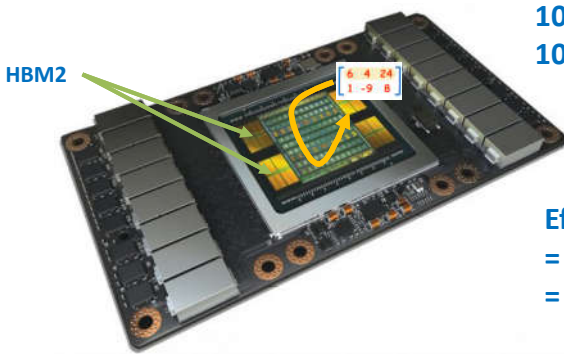
$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$

- Consider an **$n \times n$ matrix** where **32 divides n** .
- Focus on the GPU device code:
 - the host code performs typical tasks: data allocation and transfer between host and device, the launching and timing of several kernels, result validation, and the deallocation of host and device memory.
- Benchmarks illustrate this section:
 - Compare** our matrix transpose kernels against a **matrix copy kernel**,
 - for each kernel, we compute the “**effective bandwidth**”, calculated in **[GB/s]** as **twice the size of the matrix (once for reading the matrix and once for writing) divided by the time of execution**

Matrix Transpose 2

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$

- Benchmarks illustrate this section:
 - for each kernel, we compute the *effective bandwidth*, calculated in [GB/s] as **twice the size of the matrix (once for reading the matrix and once for writing) divided by the time of execution**

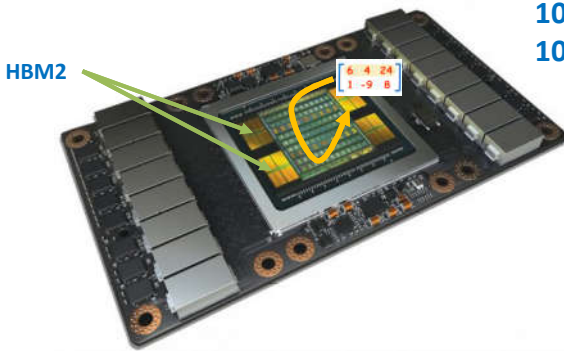


1024x1024 size float matrix
1024x1024 x4B x2 = 8MB

Effective Bandwidth
= 8MB/t_{exe}
= ?

Matrix Transpose 2

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$



1024x1024 size float matrix
1024x1024 x4B x2 = 8MB

Effective Bandwidth
= 8MB/t_{exe}
If (t_{exe}=0.02 ms) ?
= 8MB/0.02ms = 800MB/2ms = 400GB/s



Transpose 'C' code: CPU version

$O(n^2)$

```
void transpose_CPU(float in[], float out[])
{
    for(int j=0; j < N; j++)
        for(int i=0; i < N; i++)
            out[j + i*N] = in[i + j*N]; // out(j,i) = in(i,j)
}
```

<https://github.com/udacity/cs344/blob/master/Lesson%20Code%20Snippets/Lesson%205%20Code%20Snippets/transpose.cu>



Transpose 'C' code: GPU version 1

$O(n^2)$

```
// to be launched on a single thread
// transpose_serial<<<1,1>>>(d_in, d_out);
__global__ void transpose_serial(float in[], float out[])
{
    for(int j=0; j < N; j++)
        for(int i=0; i < N; i++)
            out[j + i*N] = in[i + j*N]; // out(j,i) = in(i,j)
}
```

<https://github.com/udacity/cs344/blob/master/Lesson%20Code%20Snippets/Lesson%205%20Code%20Snippets/transpose.cu>



Transpose 'C' code: GPU version 2

```
// to be launched on a single thread
// transpose_parallel_per_row<<<1,N>>>(d_in, d_out);
__global__ void
transpose_parallel_per_row(float in[], float out[])
{
    int i = threadIdx.x;

    for(int j=0; j < N; j++)
        out[j + i*N] = in[i + j*N]; // out(j,i) = in(i,j)
}
```

$O(n)$

<https://github.com/udacity/cs344/blob/master/Lesson%20Code%20Snippets/Lesson%205%20Code%20Snippets/transpose.cu>



Transpose 'C' code: GPU version 3

```
// dim3 blocks(N/K,N/K); // blocks per grid
// dim3 threads(K,K);    // threads per block
// transpose_parallel_per_element<<<blocks,threads>>>(d_in, d_out);
__global__ void
transpose_parallel_per_element(float in[], float out[])
{
    int i = blockIdx.x * K + threadIdx.x;
    int j = blockIdx.y * K + threadIdx.y;

    out[j + i*N] = in[i + j*N]; // out(j,i) = in(i,j)
}
```

$O(1)$

<https://github.com/udacity/cs344/blob/master/Lesson%20Code%20Snippets/Lesson%205%20Code%20Snippets/transpose.cu>

Deep Dive: *Performance-Aware Opt.*

Transpose 'C' code: GPU version 3


```
// dim3 blocks(N/K,N/K); // blocks per grid
// dim3 threads(K,K);    // threads per block
// transpose_parallel_per_element<<<blocks,threads>>>(d_in, d_out);
__global__ void
transpose_parallel_per_element(float in[], float out[])
{
    int i = blockIdx.x * K + threadIdx.x;
    int j = blockIdx.y * K + threadIdx.y;
    out[j + i*N] = in[i + j*N]; // out(j,i) = in(i,j)
}
```

O(1)

N = Height or width (square matrix: height == width)
K = size of a thread block in x-dimension or y-dimension

Deep Dive: *Performance-Aware Opt.*

```
ubuntu@tegra-ubuntu:~/TRANPOSE/cs344/Lesson Code Snippets/Lesson 5 Code Snippets$ ./transpose
```

transpose_serial: 963.833 ms.  ~1/74 @ N threads

Verifying transpose...Success

transpose_parallel_per_row: 13.0229 ms.

Verifying transpose...Success

transpose_parallel_per_element: 11.4738 ms.

Verifying transpose...Success

transpose_parallel_per_element_tiled 32x32: 9.39575 ms.


Verifying ...Success

transpose_parallel_per_element_tiled 16x16: 4.44258 ms.

Verifying ...Success

transpose_parallel_per_element_tiled_padded 16x16: 4.073 ms.

Verifying...Success




Matrix Transpose - CUDA

$$\begin{bmatrix} 6 & 4 & 24 \\ 1 & -9 & 8 \end{bmatrix}^T = \begin{bmatrix} 6 & 1 \\ 4 & -9 \\ 24 & 8 \end{bmatrix}$$

256 threads

- All kernels launch **thread blocks** of dimension **"32x8"**, where each block transposes (or copies) a tile of dimension 32x32. *launch blocks of 32x8 threads!*
- As such, the parameters **TILE_DIM** and **BLOCK_ROWS** are set to 32 and 8, respectively.



A simple copy kernel 1 – for Comparison

```
__global__ void copy(float *odata, float* idata)
{
    int x = blockIdx.x*TILE_DIM + threadIdx.x;
    int y = blockIdx.y*TILE_DIM + threadIdx.y;
    int width = gridDim.x * TILE_DIM;

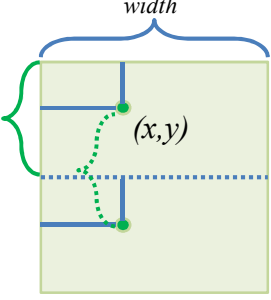
    for (int j=0; j<TILE_DIM; j+=BLOCK_ROWS) {
        odata[(y+j)*width + x] = idata[(y+j)*width + x];
    }
}
```

BLOCK_ROWS

(x, y)

width

```
j = 0
>> odata[y*width + x] = idata[y*width + x];
j = 1
>> odata[(y+1)*width + x] = idata[(y+1)*width + x];
```



<https://devblogs.nvidia.com/efficient-matrix-transpose-cuda-cc/>



A simple copy kernel 2 – for Comparison

- **odata** and **idata** are pointers to the input and output matrices,
- **width** = `gridDim.x*TILE_DIM`
- In this kernel, **xIndex** and **yIndex** are global 2D matrix indices and they used to calculate index, the 1D index used to access matrix elements.

```
__global__ void copy(float *odata, float* idata)
{
    int x = blockIdx.x*TILE_DIM + threadIdx.x;
    int y = blockIdx.y*TILE_DIM + threadIdx.y;
    int width = gridDim.x * TILE_DIM;

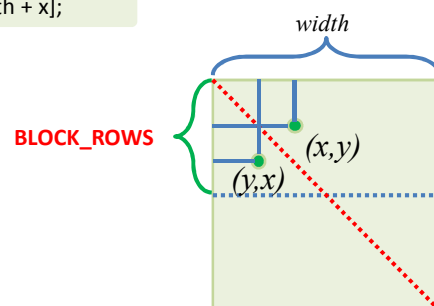

    for (int j=0; j<TILE_DIM; j+=BLOCK_ROWS) {
        odata[(y+j)*width + x] = idata[(y+j)*width + x];
    }
}
```



A naive transpose kernel

```
__global__ void transposeNaive(float *odata, float* idata)
{
    int x = blockIdx.x*TILE_DIM + threadIdx.x;
    int y = blockIdx.y*TILE_DIM + threadIdx.y;
    int width = gridDim.x*TILE_DIM;

    for (int j=0; j<TILE_DIM; j+=BLOCK_ROWS) {
        odata[x*width + (y+j)] = idata[(y+j)*width + x];
    }
}
```



<https://devblogs.nvidia.com/efficient-matrix-transpose-cuda-cc/>



Naive transpose kernel vs copy kernel

- The performance of these two kernels on a 1024x1024 matrix using a Tesla GPUs is given in the following table:

Effective Bandwidth (GB/s, ECC enabled)		
Routine	Tesla M2050	Tesla K20c
copy	105.2	136.0
transposeNaive	18.8	55.3




<https://devblogs.nvidia.com/efficient-matrix-transpose-cuda-cc/>



Problem ?

```
__global__ void transposeNaive(float *odata, float* idata)
{
    int x = blockIdx.x*TILE_DIM + threadIdx.x;
    int y = blockIdx.y*TILE_DIM + threadIdx.y;
    int width = gridDim.x*TILE_DIM;

    for (int j=0; j<TILE_DIM; j+=BLOCK_ROWS) {
        odata[x*width + (y+j)] = idata[(y+j)*width + x];
    }
}
```



(xindex, yindex): $(0, 0) \rightarrow (1, 0) \rightarrow (2, 0) \rightarrow (3, 0) \rightarrow \dots (31, 0)$
 $\rightarrow (0, 1) \rightarrow (1, 1) \rightarrow (2, 1) \rightarrow (3, 1) \rightarrow \dots (31, 1)$
 $\rightarrow \dots$

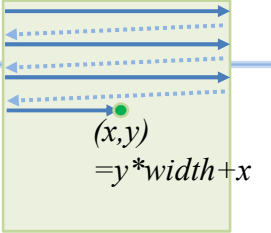
Problem ?

```

__global__ void transposeNaive(float *odata, float* idata)
{
    int x = blockIdx.x*TILE_DIM + threadIdx.x;
    int y = blockIdx.y*TILE_DIM + threadIdx.y;
    int width = gridDim.x*TILE_DIM;

    for (int j=0; j<TILE_DIM; j+=BLOCK_ROWS) {
        odata[x*width + (y+j)] = idata[(y+j)*width + x];
    }
}

```



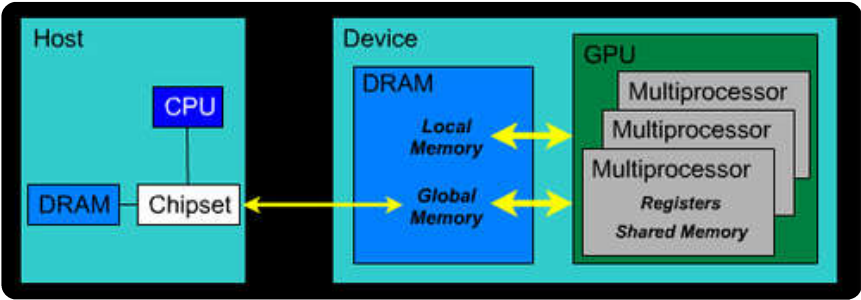
(x,y)
=y*width+x

(xindex, yindex): (0, 0) → (1, 0) → (2, 0) → (3, 0) → ... (31, 0)
 → (0, 1) → (1, 1) → (2, 1) → (3, 1) → ... (31, 1)
 → ...

One Transaction .vs. 32 Transactions !

Coalesced Transpose 1


- Because **device memory [GPU DDR Memory]** has a much higher latency and lower bandwidth than **on-chip memory [shared memory]**, special attention must be paid to: **how global memory accesses are performed?**





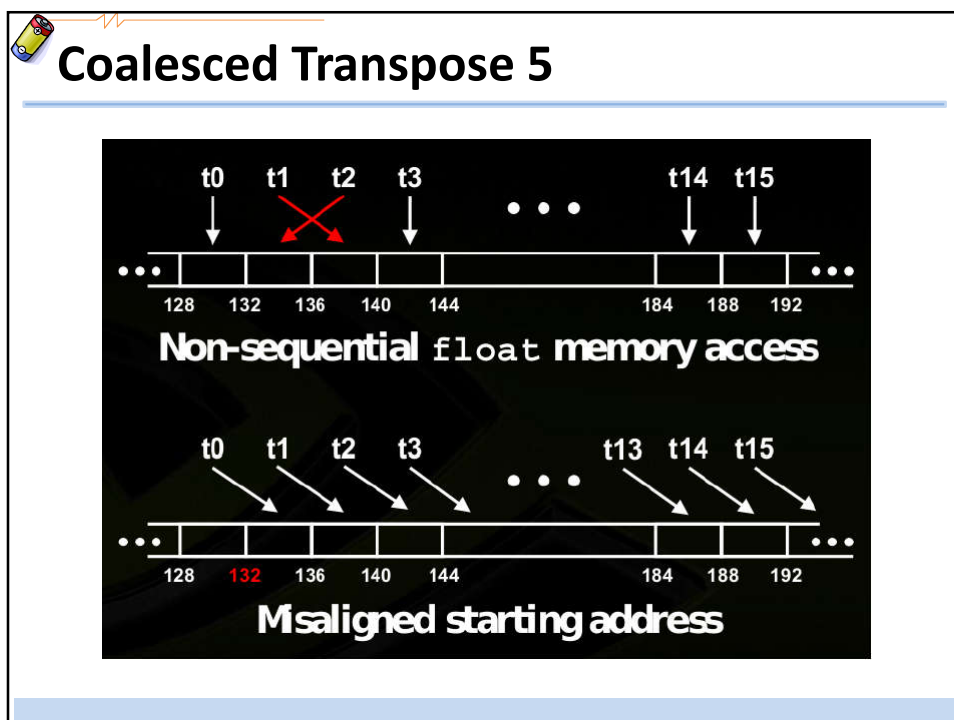
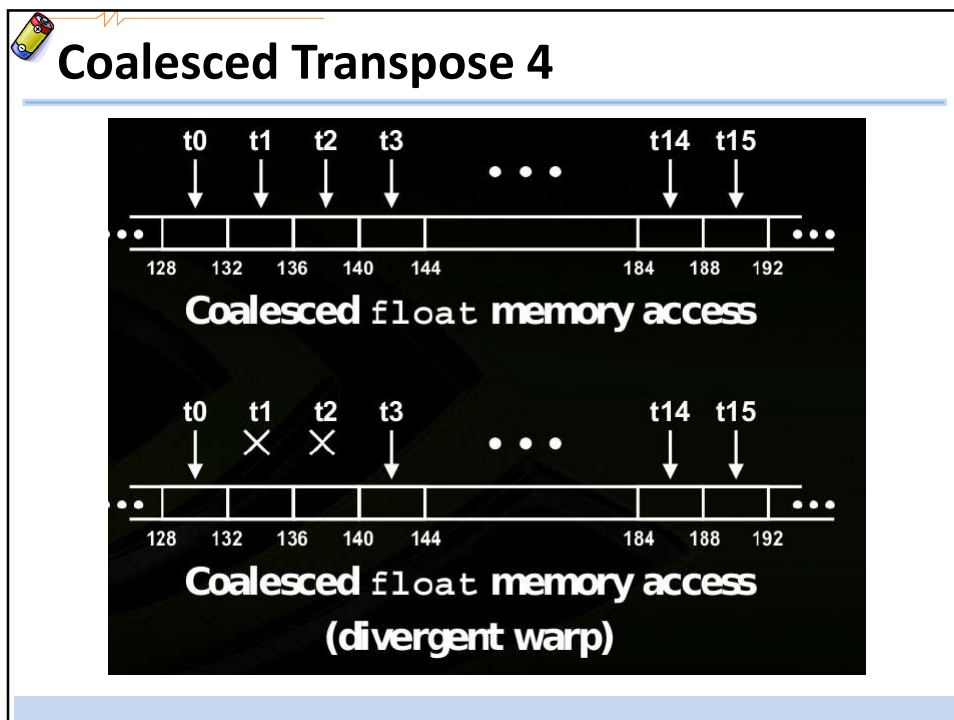
Coalesced Transpose 2

- The **simultaneous global memory accesses** by each thread of a during the execution of a single read or write instruction will be **coalesced** into a single access if:
 - The size of the memory element accessed by each thread is either 4, 8, or 16 bytes.
 - The elements form a **contiguous block of memory**.
 - The i-th element is accessed by the i-th thread in the warp.



Coalesced Transpose 3

- The simultaneous global memory accesses by each thread of a during the execution of a single read or write instruction will be **coalesced** into a single access if:
 - The size of the memory element accessed by each thread is either 4, 8, or 16 bytes.
 - The elements form a **contiguous block of memory**.
 - The i-th element is accessed by the i-th thread in the warp.
- Last two requirements can be relaxed** (compiler optimization) with compute capabilities of 1.2.
- Coalescing happens** even if some threads do not access memory (**divergent warp**)





Coalesced Transpose 6

- Basically, “all loads from idata are coalesced”.
- Coalescing behavior differs between the simple copy and naïve transpose kernels when **writing to odata**.
 - Simple copy – coalesced
 - Naïve Transpose – non-coalesced



Coalesced Transpose 7

- The way to avoid **uncoalesced** global memory access is
 - to read the data into **shared memory** and,
 - have each warp access noncontiguous locations in shared memory in order to write contiguous data to odata.
- There is no performance penalty for noncontiguous access patterns in shared memory as there is in global memory.
- a **__syncthreads()** call is required to ensure that all reads from idata to shared memory have completed before writes from shared memory to odata.



Coalesced Transpose 8

```

__global__ void transposeCoalesced(float *odata, const float *idata)
{
    __shared__ float tile[TILE_DIM][TILE_DIM];

    int x = blockIdx.x * TILE_DIM + threadIdx.x;
    int y = blockIdx.y * TILE_DIM + threadIdx.y;
    int width = gridDim.x * TILE_DIM;

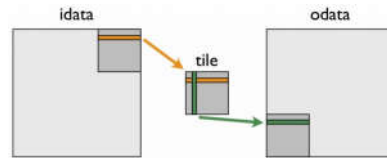
    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS)
        tile[threadIdx.y+j][threadIdx.x] = idata[(y+j)*width + x];

    __syncthreads();

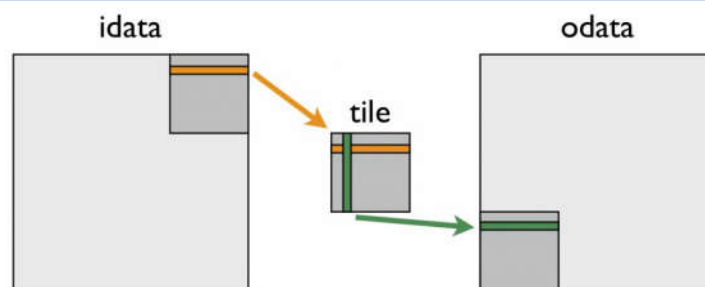
    x = blockIdx.y * TILE_DIM + threadIdx.x; // transpose block offset
    y = blockIdx.x * TILE_DIM + threadIdx.y;

    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS)
        odata[(y+j)*width + x] = tile[threadIdx.x][threadIdx.y + j];
}

```



Coalesced Transpose 9



- a warp of threads reads contiguous data from idata into rows of the shared memory tile.
- After recalculating the array indices, **a column of the shared memory tile is written to contiguous addresses in odata.**
- Because threads write different data to odata than they read from idata, we must use a block-wise barrier synchronization **__syncthreads()**.



Coalesced Transpose 10

Effective Bandwidth (GB/s, ECC enabled)		
Routine	Tesla M2050	Tesla K20c
copy	105.2	136.0
copySharedMem	104.6	152.3
transposeNaive	18.8	55.3
transposeCoalesced	51.3	97.6

- There is a dramatic increase in effective bandwidth of the coalesced transpose over the naive transpose, **but there still remains a large performance gap** between the coalesced transpose and the copy:
 - One possible cause of this performance gap could be the **synchronization barrier** required in the coalesced transpose.
 - This can be easily assessed using the following copy kernel which utilizes shared memory and contains a `__syncthreads()` call.



Coalesced Transpose 11

```

__global__ void copySharedMem(float *odata, const float *idata)
{
    __shared__ float tile[TILE_DIM * TILE_DIM];

    int x = blockIdx.x * TILE_DIM + threadIdx.x;
    int y = blockIdx.y * TILE_DIM + threadIdx.y;
    int width = gridDim.x * TILE_DIM;

    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS)
        tile[(threadIdx.y+j)*TILE_DIM + threadIdx.x] = idata[(y+j)*width + x];

    __syncthreads();

    for (int j = 0; j < TILE_DIM; j += BLOCK_ROWS)
        odata[(y+j)*width + x] = tile[(threadIdx.y+j)*TILE_DIM + threadIdx.x];
}


```

<https://devblogs.nvidia.com/efficient-matrix-transpose-cuda-cc/>

Coalesced Transpose 12

Effective Bandwidth (GB/s, ECC enabled)		
Routine	Tesla M2050	Tesla K20c
copy	105.2	136.0
copySharedMem	104.6	152.3
transposeNaive	18.8	55.3
transposeCoalesced	51.3	97.6

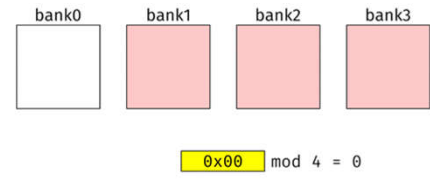
- The shared memory copy results seem to suggest that **the use of shared memory with a synchronization barrier has little effect on the performance.**



Shared memory bank conflicts 1

- Shared memory is divided into 32 equally-sized memory modules, called **banks**, which are organized such that successive 32-bit words are assigned to successive banks.

Write



- These banks can be accessed simultaneously, and to achieve maximum bandwidth to and from shared memory, the **[threads in a warp should access shared memory associated with different banks]**.

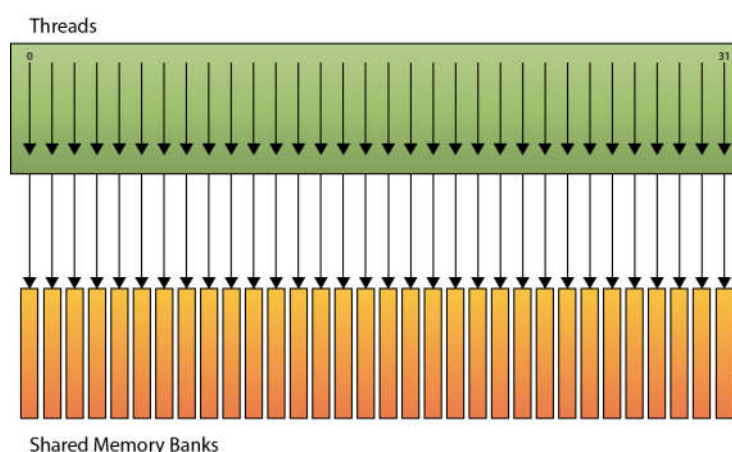


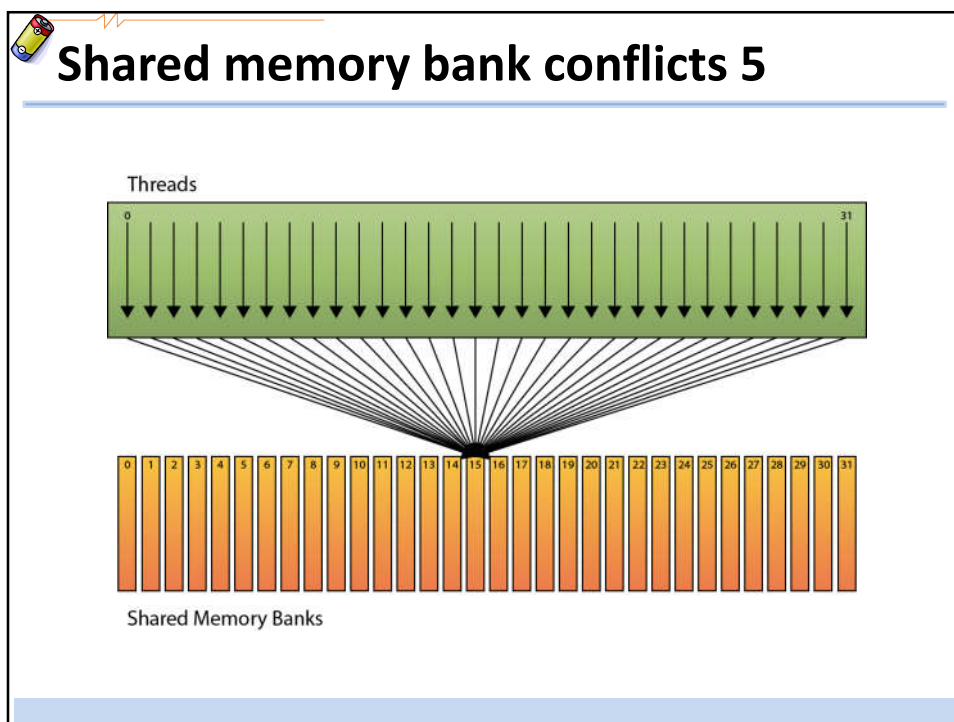
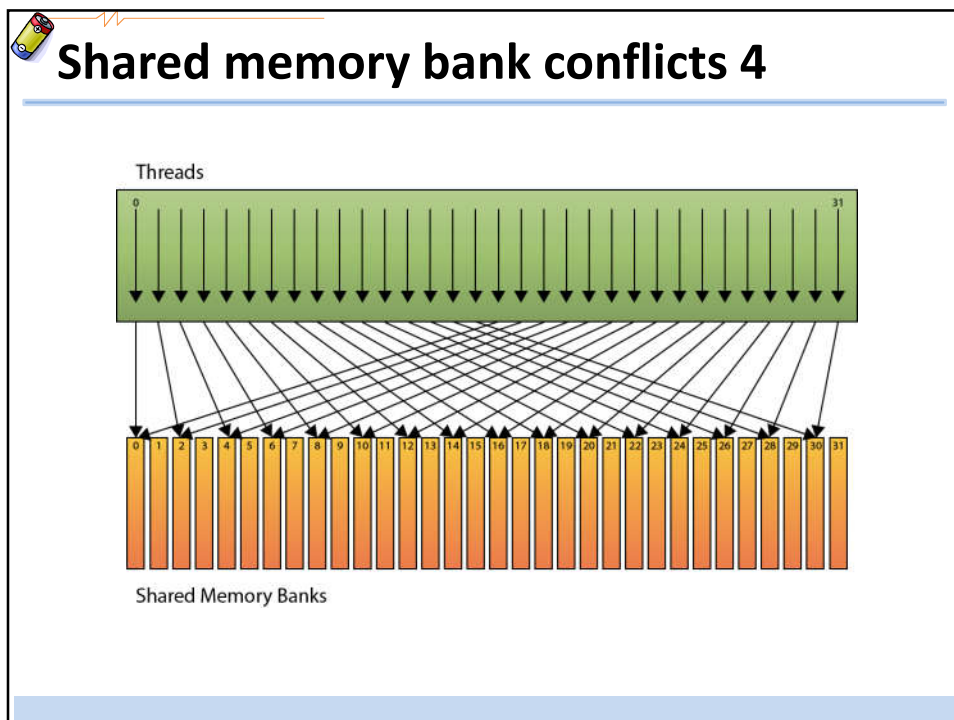
Shared memory bank conflicts 2

- These banks can be accessed simultaneously, and to achieve maximum bandwidth to / from shared memory the **threads in a warp should access a shared memory associated with different banks**.
- The **exception to this rule** is when **all threads in a warp read the same shared memory address**, which results in a **broadcast** where the data at that address is sent to all threads of the half warp in one transaction.



Shared memory bank conflicts 3







Shared memory bank conflicts 6

- The coalesced transpose uses a 32x32 shared memory array of floats.
- For a shared memory tile of 32×32 elements, **all elements in a column of data map to the same shared memory bank**
 - Resulting in a worst-case scenario for memory bank conflicts: reading a column of data results in a **32-way bank conflict**.
- A simple way to avoid this conflict is to **pad the shared memory array by one column**:

```
__shared__ float tile[TILE_DIM][TILE_DIM+1];
```



Shared memory bank conflicts 7

- A simple way to avoid this conflict is to **pad the shared memory array by one column**:

```
__shared__ float tile[TILE_DIM][TILE_DIM+1];
```

bnk1	bnk2	bnk3	bnk4
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16




bnk1	bnk2	bnk3	bnk4
1	2	3	4
p	5	6	7
8	p	9	10
11	12	p	13
14	15	16	p

tile[tid % 4]:

→ 1, 5, 9, 13 threads access a bank1

Shared memory bank conflicts 8

Effective Bandwidth (GB/s, ECC enabled)		
Routine	Tesla M2050	Tesla K20c
copy	105.2	136.0
copySharedMem	104.6	152.3
transposeNaive	18.8	55.3
transposeCoalesced	51.3	97.6
transposeNoBankConflicts	99.5	144.3

 SATISFIED!

<https://devblogs.nvidia.com/efficient-matrix-transpose-cuda-cc/>

Granularity of Parallelism

- **Size of a Tile ?**
 - We do test with a block of 32x8 threads with config. of “(32,8)”
 - What about **32x32** ?
 - “1024 threads wait at a barrier”
 - High Parallelism (?) but high synchronization overhead
 - What about **16x16** ?
 - “256 threads wait at a barrier”
 - Lower Parallelism (?) but lower synchronization overhead

Minimize timing waiting at a barrier !

<https://github.com/jeonggunlee/cs344/blob/master/Lesson%20Code%20Snippets/Lesson%205%20Code%20Snippets/transpose.cu>



Conclusion - Transpose

- Understand CUDA performance characteristics
 - **Memory coalescing**
 - **Bank conflicts**
 - **Granularity of parallelism**
- Use peak performance metrics to guide optimization